

My notes on Linux systems

Thanasisn github.com/thanasisn

2025-05-03

Contents

	4
Intro	5
0.1 Backup your precious data	5
0.2 Distro hopping	6
0.3 Design and reasoning	6
1 System Install	8
1.1 The .iso	8
1.2 Partitioning and file system.	8
1.3 Install the basics	8
2 System configuration	9
2.1 System mail over gmail with postfix	9
2.2 collectd	10
2.3 ntp	11
2.4 VPN network	11
2.5 Setup .ssh/config	11
2.6 Set ssh keys	11
2.7 Keep your files and config in sync	12
2.8 Install and configure zsh	12
3 System Backup	12
3.1 Software	13
3.2 Current scheme	13
4 Random stuff, tricks and tools	14
4.1 ddresque	14
4.2 dd and gzip img file	14
4.3 Polar Pro Trainer 5 and IRDA	15
4.4 OCR	15
4.5 Spell check word list	16
4.6 Mount physical HDD on a VirtualBox machine	17
4.7 Ignore lid close on laptops	17
5 Software Selection Reasoning	17
5.1 Mail client	17
5.2 File browser	18
5.3 General text editing	18

6	Commands I forget	18
6.1	Network monitoring	18
6.2	System monitoring	18
6.3	Testing	18
7	btrfs	19
7.1	Break a mirror	19
7.2	Rebalance to reduce unallocated space	19
7.3	Just format two disks into raid1 with:	20
7.4	Add new disk to raid1	20
8	zfs	20
8.1	from a pool history	21

[pdf](#) file

This work is licensed under
Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International
creativecommons.org/licenses/by-nc-nd/4.0/

For correction or contributions contact ThanasisN at github.com/thanasisn.

Notice: These notes are for my own use. They are opinionated and may contain any kind of errors.

Intro

Here I am collecting my notes about deploying GNU/Linux systems. **This is mainly for my reference**, for setting up a system with my preferred setup and functionality. In general, I am guided by minimal and functional choices. I will go for simple and minimal solutions, but will not avoid bloating it, if it is easier, at least for a first approach solution. Many times it is easier to start with a GUI for new stuff and ditch it later, as you learn more.

Most stuff are tested and used on “Debian”, but I have an eye for the “NixOS” way.

This is not trying to be a step by step guide. You can google a thousand tutorials for each subject.

If you find it useful, you can contribute to my coffee addiction, with a donation here:



A lot of my code lives here: github.com/thanasisn, along with Linux tools and research stuff.

As Bill† used to say ...

Before touching the keyboard, make a plan and design, what you are going to do with it.

0.1 Backup your precious data

0.1.1 First approach

Don't back up anything. You only live once and your data won't matter to the Universe.

0.1.2 Another approach

- **Backup everything**, your files, your responsibility.
- Verify your backups still work, regularly.

- Get things that you may need from all of your systems.
 - /etc config files
 - * /etc/tinc
 - * /etc/collectd
 - * /var/lib/ganglia
 - * /var/lib/collectd
 - * /var/spool/cron
 - ssh keys from users
 - user profile files .bashrc .history etc.
 - lists of installed packages
 - other software installed manual
 - things that are not in your \$HOME
 - files from ‘root’ user account
 - list of libraries from your programming languages
 -

Your backup should be automated, and the configuration collection scripted ([example](#)).

0.2 Distro hopping

Do it, and get over it. Use a test machine or a virtual machine to fool around, keep your main machine always available.

- Get you favorite distro. Keep it simple and robust with “Debian or Arch”.
- Prefer net install media. We start with a minimal install and will build up.
- Check if you need firmware. After formatting is not easy to get to the net for answers and software.

0.3 Design and reasoning

0.3.1 For building a system, I want:

Fast setup: I tend to set up and format secondary machines regularly.

Reproducible environments: I like to have the same functionality on any machine I use.

Common things: Security, syncing, settings among the hosts.

Evolution: The ability, and the opportunity to improve my setups.

0.3.2 My systems usually are:

Main personal system: A laptop at home, everything starts from here.

Desktop system for work: Work horse and backup.

Grid/cluster computing account: Configs and programs copied there to increase functionality.

Server systems for work: Storage, compute nodes and more.

Secondary personal system: A laptop for working outside with extra security.

Home server: Mainly backup storage and cloud uploader.

Secondary and old machines: Media players, test machines ...

So, these notes help me, and reminds me stuff to keep all this running and happening.

0.3.3 Current system goal

Root on an btrfs with raid 1: Two disk for redundancy end error correction checking against bit rot. Easier installation on Debian.

Home partition on encrypted zfs raid 1: Redundancy, error checking and safety. Not so safe as with full disk encryption but good enough. Now I am using a 'gocryptfs' filesystem on top of a btrfs partition.

Auto login on boot without mounting users encrypted home: So I can ssh from my other machines and act as a honeypot in case of losing it.

Manual mounting/unmounting of users encrypted home: So, I have control of the data exposure. Use a different .ssh/key for the two states, so other application can work normally while encrypted (sync, etc).

0.3.4 Security and privacy

You should really consider some things about security and privacy.

Check that: wiki.debian.org/SetupGuides/SecurePersonalComputer

Also you should use email aliases, and pgp encryption, especially when using gmail and similar

providers.

1 System Install

1.1 The .iso

- Grub the net installer for your favorite distro.
- Check for firmware dependencies.
- Use the magic of ventoy.net

1.2 Partitioning and file system.

Think about partitioning and file system you are going to use.

- is a multi-boot?
- want encrypted partitions?
- keep server functioning even with encrypted partition?

It is easier to set it or prepare it, at system install, especially if you want to make something complex later.

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPPOINT
sda	8:0	0	465.8G	0	disk	
sda1	8:1	0	953M	0	part	/boot
sda2	8:2	0	1K	0	part	
sda3	8:3	0	421.1G	0	part	
crypthome	254:0	0	421G	0	crypt	/home/athan
sda5	8:5	0	14G	0	part	/
sda6	8:6	0	14G	0	part	/home/arch
sda7	8:7	0	1.9G	0	part	[SWAP]

Here is my old scheme. Dual boot with a spare partition and an encrypted home partition, all on btrfs file system.

1.3 Install the basics

Install the most basic options your distro provides.

- No desktop environment
- No GUI apps

Start just with a terminal and an ssh server, and get only the things you need. Usually I have

a minimal DE like `lxde` for fail back or guest users, I use `dwm` as window manager, and mix GUI apps from different desktops. Of course the looks of the windows become inconsistent, but you can configure your way out of this if you want.

I try to purge any program I don't need or stopped using.

2 System configuration

Mostly based on `debian`.

2.1 System mail over gmail with postfix

i.e. I couldn't make `exif4` work for me

2.1.1 Prepare

Issue a unique password from <https://myaccount.google.com/apppasswords> to use for each machine.

2.1.2 Install

```
sudo apt-get install libsasl2-modules postfix
```

2.1.3 Configure

Choose:

Internet Site

Edit:

```
sudo vim /etc/postfix/main.cf
```

```
. . . . .  
## add  
smtp_use_tls = yes  
smtp_sasl_auth_enable = yes  
smtp_sasl_security_options =  
smtp_sasl_password_maps = hash:/etc/postfix/sasl_passwd  
smtp_tls_CAfile = /etc/ssl/certs/ca-certificates.crt  
. . . . .  
## edit
```

```
relayhost = [smtp.gmail.com]:587  
.  
.  
.  
.  
.
```

Create password file:

```
sudo vim /etc/postfix/sasl_passwd
```

```
[smtp.gmail.com]:587    here_is_your_mail@gmail.com:pasw0rd_from_google_app_unique_passwords
```

Prepare password for use:

```
sudo postmap /etc/postfix/sasl_passwd
```

Edit system aliases:

```
sudo vim /etc/aliases
```

Add to whom the mail is going.

Restart services:

```
sudo newaliases  
sudo systemctl restart postfix
```

2.2 collectd

I use this to monitor multiple machines, and have graphs of all on my conky.

2.2.1 Install

```
apt-get install lm-sensors hddtemp collectd ntp
```

2.2.2 Configure

```
sudo /lib/systemd/systemd-sysv-install enable hddtemp  
sudo sensors-detect  
sudo vim etc/collectd.conf
```

2.2.3 Starting the daemon

You need to restart the daemon.

```
sudo /etc/init.d/collectd restart
```

or

```
sudo systemctl enable collectd.service  
sudo systemctl restart collectd.service
```

For hddtemp have to enable daemon and start service

```
sudo vim /etc/default/hddtemp  
sudo systemctl enable hddtemp.service
```

2.3 ntp

Keep your clock in sync

```
## enable  
sudo systemctl enable ntp  
## start running  
sudo systemctl restart ntp  
## check  
ntpq -p
```

2.4 VPN network

If you have more than one machine, **network all devices in a private VPN**. I use `tinc`.

It is easier if you have available a machine with static IP. Otherwise, you can use dynamic DNS services to do that.

Having all your devices able to talk to each other makes everything easier. You can then, easily setup ssh, backup, nfs, file sync and much more. You can access hosts even with dynamic IP and behind firewalls and NAT.

2.5 Setup `.ssh/config`

Use the `.ssh/config` file to add all your hosts and all their IPs public or private. This way, you don't have to write the addresses or give the identity file each time. And also, this will work for many other applications that use ssh configuration.

2.6 Set ssh keys

```
ssh-keygen -t dsa
chmod 600 .ssh/id_dsa
scp .ssh/id_dsa.pub dude@remote.edu:
cat id_dsa.pub >> .ssh/authorized_keys
chmod 700 .ssh
```

2.7 Keep your files and config in sync

Using unison to sync config files, so all machines feels the same.

I select with files are synced based on each machine's role, work, personal, testing...

That doesn't count as backup. There are cases that I have messed thing up with it.

2.8 Install and configure zsh

```
sudo apt-get install zsh
```

Get oh-my-zsh from its repo.

github.com/robbyrussell/oh-my-zsh

Change default shell

```
chsh -s /usr/bin/zsh
```

2.8.1 change shell

```
chsh -s /usr/bin/zsh
```

3 System Backup

Your backup should be automated! human errors and reliability is a constant problem for computers. If you think, copying data by hand to a USB drive is a backup, you are wrong. **COPYING DATA BY HAND TO A USB DRIVE IS POINTLESS.** It will be lost, neglected and fail at some point, sooner than you expect.

Testing your backup should also be automated! You have to know it works, not assume it.

Constantly log and report on the backup status! You want to know if your backup has failed, now. Not three months before you need it.

I have a script that shows on conky the status of each backup, and counts how long before was the last successful backup, last warning and last fail. Also monitor the size of each as an indication of things going good. Color coded on conky so if I see

- white, everything is normal don't bother to look
- green, something is out of the ordinary but still good
- yellow, a problem has emerged, may be in danger, take action
- red, this is off, consider it not safe

3.1 Software

Backup programs, I have extensively used for years on multiple machines at the past:

- `deja-dup` too long ago to remember
- `backintime` elementary but good
- `bacula` excellent

Current program `borg`. Deduplication, encryption, compression, snapshots everything I need and simple. Upload the whole backup data to the cloud with `rclone`.

Some custom bash scripts on cron do everything, backup, prune, test, upload, report.

3.2 Current scheme

- Local backup
 - One `raid 1` array for high frequency sort-term retention (`borg` stable version)
 - One `btrfs` `raid1` array for low frequency midterm retention (`borg` testing version)
 - An external drive for occasional easy portable backup (`borg` testing version)
- Remote backup
 - One on a private remote machine (`borg` stable version)
- Cloud
 - Sync a backup copy from local to google drives (`rclone`)
 - Sync a new backup copy every month from remote to google drive (`rclone`)

This is for backing up my personal main machine. Occasionally, I have a local backup with `borg` for some of the secondary machines. These include config files or a lot of data, that can be recovered with other means but will take time and bandwidth.

4 Random stuff, tricks and tools

4.1 ddrescue

Rescue failing partitions to a new disk or to an image file with ddrescue.

4.1.1 First pass

Do a “fast” copy ignoring hard problems.

```
ddrescue -v -f -n -a 5120000 /dev/sda /dev/sdc logfile.log
```

- -n keeps ddrescue from “scraping”, that puts a heavy load on the drive trying to rescue the most difficult parts
- -a 5120000 it will skip any area that cannot be copied at a speed of at least 5 megabytes per second
- logfile.log a file to use, so you can start over from where you left off rather than starting from start
- -f force overwrite of outfile, needed when outfile is not a regular file, but a device or partition
- -v verbose

4.1.2 Second pass

Do a more intensive try to copy.

```
ddrescue -v -f -d -r 3 /dev/sda /dev/sdc logfile.log
```

- -d direct disk access
- -r 3 read unrecoverable sectors 3 times
- -f force overwrite of outfile, needed when outfile is not a regular file, but a device or partition
- -v verbose

4.2 dd and gzip img file

4.2.1 dd disk to compressed image

```
dd if=/dev/sdx | gzip > /path/to/image.gz
```

4.2.2 dd compressed image to disk

```
gzip -dc /path/to/image.gz | dd of=/dev/sdx
```

4.3 Polar Pro Trainer 5 and IRDA

4.3.1 Polar Pro Trainer 5

This is a 32 bit program, so we need wine32.

```
## install wine
sudo dpkg --add-architecture i386 && sudo apt update
sudo apt-get install wine32
## setup wine
export WINEARCH=win32
export WINEPREFIX=~/.wine32
winecfg
## run program setup
wine setup.exe
```

You can update the program with overwriting newer files.

4.3.2 Setup IRDA probe

```
sudo apt-get install irda-utils
```

For configuring ma-620 there is a bash script: [ma-620_setup.sh](#).

For polar USB dongle setup

```
alias irda0 mcs7780
```

4.3.2.1 /etc/modprobe.d/irda-utils.conf

```
ENABLE="true"
DEVICE="irda0"
MAX_BAUD_RATE="9600"
```

4.3.2.2 /etc/default/irda-utils

4.4 OCR

4.4.1 ocrfeeder

- using tesseract

4.4.2 gimagereader

- easier and prettier than ocrfeeder
- using tesseract

4.4.3 ocrmypdf

- nice
- CLI

```
ocrmypdf -l eng+ell --rotate-pages --deskew
ocrmypdf -l eng+ell --rotate-pages --deskew --clean --remove-background
```

Usually I want to use this:

```
ocrmypdf --jpeg-quality 100 --png-quality 100 --optimize 0 -l eng+ell --rotate-pages --deskew
```

Part of man page...

```
--optimize {0,1,2,3}
    Control how PDF is optimized after processing:
    0 - do not optimize;
    1 - do safe, lossless optimizations (default);
    2 - do some lossy optimizations;
    3 - do aggressive lossy optimizations (including lossy JBIG2)

--deskew
    Deskew each page before performing OCR

--clean
    Clean pages from scanning artifacts before performing OCR,
    and send the cleaned page to OCR, but do not include the
    cleaned page in the output
```

4.5 Spell check word list

I hard link all the files of the same language so

- All are updated instantly with new words
- I can remove words or edit all file together
- Use common dict for all programs

Found linked files with: `find -L ./ -samefile .vim/spell/el.utf-8.add`


```

./.config/enchant/el.dic
./.config/enchant/el_GR.dic
./.hunspell_el_GR
./.vim/bundle/vim-grammarous/misc/LanguageTool-5.5/org/linguist/grammarous/grammarous/el/hunspell/spell
./.vim/spell/el.utf-8.add

./.config/enchant/en.dic
./.config/enchant/en_GB.dic
./.config/enchant/en_US.dic
./.hunspell_en_US
./.vim/bundle/vim-grammarous/misc/LanguageTool-5.5/org/linguist/grammarous/grammarous/en/hunspell/spell
./.vim/spell/en.utf-8.add

```

- Some programs I use like Rstudio and textstudio are missing.
- vim doesn't add words to the correct language.

4.6 Mount physical HDD on a VirtualBox machine

You can do a fast OS installation through a virtual machine on a physical hard drive.

```

sudo VBoxManage internalcommands createrawvmdk -filename "/path/to/file.vmdk" -rawdisk /dev/sda
sudo chmod 777 /dev/sda
sudo chmod 777 /path/to/file.vmdk

```

4.7 Ignore lid close on laptops

Edit: /etc/systemd/logind.conf

```

.....
HandleLidSwitch=ignore
.....

```

```

sudo service systemd-logind restart

```

5 Software Selection Reasoning

A reminder to myself of which programs I have tried, and why I use them.

5.1 Mail client

Tested: thunderbird, kontact and evolution.

I have used them all for a long time with multiple accounts. I end up with evolution. As long

as I have seahorse unlocked everything works fine. The other two are good programs but the synchronization of options across machines is not so straightforward.

5.2 File browser

Tested: a bunch of the most prominent.

End up working with terminal, caja and lf, and occasional dolphin and mc. All are mapped to keyboard keys with sxhkd, so I start with the most appropriate, for what I have in mind.

5.3 General text editing

Nothing beats vim. Learn now, thank me later.

Regularly, I also use kate and sometimes geany.

6 Commands I forget

Less often used command that are very handy, but I tend to forget due to infrequent usage.

6.1 Network monitoring

nethogs per program network bandwidth monitor

iftop -i any: per connection bandwidth and volume monitor

6.2 System monitoring

htop general system monitoring

bpytop general system monitoring/manager

glances general system monitoring

6.3 Testing

- `systemctl enable ssh.service`
- `systemctl start ssh.service`
- `PasswordAuthentication yes`
- `sudo adduser bob`
- `sudo passwd`
- `vim /etc/hostname`
- `vim /etc/hosts`
- `visudo`
- `sudo vim /etc/lightdm/lightdm.conf`
- `sudo apt install tinc`

7 btrfs

I use btrfs mostly to it's flexibility. Mostly when I foresee the need to make changes on the partitions or the disks.

7.1 Break a mirror

The functionality of btrfs continues to evolve over time. The `delete` command is now an alias for the `remove` command which produces different results. Today, you would run the following commands:

```
btrfs balance start -f -sconvert=single -mconvert=single -dconvert=single <mount>
btrfs device remove <drive> <mount>
```

The first command converts all data from a mirrored setup to a single-copy setup. This effectively makes the RAID1 into a JBOD setup. The `-f` option is required to tell the filesystem to *really* reduce the resiliency of the data.

Once this completes, the second command removes the device from the JBOD. The filesystem will move any data from the removed device to the other device.

7.2 Rebalance to reduce unallocated space

Avoid full balance

```
sudo btrfs balalance start -dusage=40 /
sudo btrfs balalance start -dusage=25 /media/maindata
```

Use for new disk on raid

```
sudo btrfs balance start /
sudo btrfs balance start /media/maindata
```

Watch the balance progress

```
sudo watch "btrfs filesystem usage -T /; btrfs balance status /"
sudo watch "btrfs filesystem usage -T /media/maindata; btrfs balance status /media/maindata"
```

Info on the filesystem

```
sudo btrfs device stats /dev/sda4
```

```
sudo btrfs device stats /dev/sda5
```

7.3 Just format two disks into raid1 with:

```
sudo mkfs.btrfs -m raid1 -d raid1 /dev/sdb /dev/sdc
```

7.4 Add new disk to raid1

List existing file systems

```
sudo btrfs filesystem show
```

```
Label: 'barelfs'  uuid: bd59a77e-a39c-4dad-885c-d956bc67d12b
```

```
    Total devices 1 FS bytes used 1.16TiB
```

```
    devid    1 size 1.77TiB used 1.16TiB path /dev/sdd2
```

Add disk to the filesystem

```
sudo btrfs device add -f /dev/sde2 /media/barel/
```

Start balancing to actually create the raid mirror

```
sudo btrfs balance start -dconvert=raid1 -mconvert=raid1 /media/barel/
```

8 zfs

I use ZFS mostly for the home folder to use native encryption

```
zpool create -o atime=off -o encryption=on -o keyformat=passphrase poolname drive
```

Use native encryption

```
zpool create -O atime=off -O encryption=on -O keyformat=passphrase poolname mirror drive1 drive2
```

```
zpool attach poolname existinghdd newhdd
```

```
zpool status poolname
```

```
sudo zpool create -f test mirror /home/athan/ZHOST/4G.img /home/athan/ZHOST/3G.img
```

8.1 from a pool history

```
zpool create maindata raidz1 /dev/sdb3 /dev/sda2
zfs set atime=off maindata
zfs create -o encryption=on -o keyformat=passphrase maindata/home_store
zfs set mountpoint=/home/ maindata/home_store
zpool import -c /etc/zfs/zpool.cache -aN
zfs load-key -r maindata/home_store
zfs create -o mountpoint=/nix maindata/nix_store
zfs set atime=off maindata
zpool set autotrim=on maindata
zfs set xattr=sa maindata
zpool import -c /etc/zfs/zpool.cache -aN
zfs load-key -r maindata/home_store
```